

```
!pip install linearmodels
```

```
Requirement already satisfied: linearmodels in /usr/local/lib/python3.12/dist-packages (7.0)
Requirement already satisfied: numpy<3,>=1.22.3 in /usr/local/lib/python3.12/dist-packages (from linearmodels) (2
Requirement already satisfied: pandas>=1.4.0 in /usr/local/lib/python3.12/dist-packages (from linearmodels) (2.2.
Requirement already satisfied: scipy>=1.8.0 in /usr/local/lib/python3.12/dist-packages (from linearmodels) (1.16.
Requirement already satisfied: statsmodels>=0.13.0 in /usr/local/lib/python3.12/dist-packages (from linearmodels)
Requirement already satisfied: mypy_extensions>=0.4 in /usr/local/lib/python3.12/dist-packages (from linearmodels)
Requirement already satisfied: pyhdfs>=0.1 in /usr/local/lib/python3.12/dist-packages (from linearmodels) (0.2.0)
Requirement already satisfied: formulaic>=1.2.1 in /usr/local/lib/python3.12/dist-packages (from linearmodels) (1
Requirement already satisfied: interface-meta>=1.2.0 in /usr/local/lib/python3.12/dist-packages (from formulaic>
Requirement already satisfied: narwhals>=1.17 in /usr/local/lib/python3.12/dist-packages (from formulaic>=1.2.1->
Requirement already satisfied: typing-extensions>=4.2.0 in /usr/local/lib/python3.12/dist-packages (from formulai
Requirement already satisfied: wrapt>=1.0 in /usr/local/lib/python3.12/dist-packages (from formulaic>=1.2.1->line
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.12/dist-packages (from pandas>=1.
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.12/dist-packages (from pandas>=1.4.0->linea
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.12/dist-packages (from pandas>=1.4.0->lin
Requirement already satisfied: patsy>=0.5.6 in /usr/local/lib/python3.12/dist-packages (from statsmodels>=0.13.0-
Requirement already satisfied: packaging>=21.3 in /usr/local/lib/python3.12/dist-packages (from statsmodels>=0.13
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.12/dist-packages (from python-dateutil>=2.8.2->
```

```
from google.colab import files
import pandas as pd
import numpy as np
import io

uploaded = files.upload()
df = pd.read_csv(io.BytesIO(list(uploaded.values())[0]))
df.head()
```

Choose files macro_panel_updated.csv
macro_panel_updated.csv(text/csv) - 151209 bytes, last modified: 19/11/2025 - 100% done
Saving macro_panel_updated.csv to macro_panel_updated (3).csv

	country	year	quarter	npl_ratio	euribor_3m	gdp_growth	hicp_yoy	unemployment	inflation_yoy	grid
0	AT	2000	1	NaN	3.542300	4.600445	0.0	3.8	1.733333	more
1	AT	2000	2	NaN	4.263000	4.600445	0.0	3.8	1.933333	
2	AT	2000	3	NaN	4.737600	4.600445	0.0	3.8	2.033333	
3	AT	2000	4	NaN	5.024167	4.600445	0.0	3.8	2.100000	
4	BE	2000	1	NaN	3.542300	5.806174	0.0	7.1	1.633333	

Next steps: [Generate code with df](#) [New interactive sheet](#)

```
# Standardize
df.columns = df.columns.str.strip().str.lower()

# DROP the wrong inflation column
if "hicp_yoy" in df.columns:
    df = df.drop(columns=["hicp_yoy"])

# Keep only the correct inflation column
print(df.columns.tolist())

['country', 'year', 'quarter', 'npl_ratio', 'euribor_3m', 'gdp_growth', 'unemployment', 'inflation_yoy']
```

```
df["time_id"] = df["year"] * 4 + df["quarter"]

num_cols = ["year", "quarter", "time_id", "npl_ratio",
            "gdp_growth", "unemployment", "inflation_yoy", "euribor_3m"]

for c in num_cols:
    df[c] = pd.to_numeric(df[c], errors="coerce")
```

```
df = df.dropna(subset=["npl_ratio"])
df = df.sort_values(["country", "time_id"]).reset_index(drop=True)

df.head()
```

	country	year	quarter	npl_ratio	euribor_3m	gdp_growth	unemployment	inflation_yoy	time_id	grid icon
0	AT	2006	1	2.739429	2.611567	5.380802	5.7	1.433333	8025	info icon
1	AT	2006	2	2.739429	2.889500	5.380802	5.7	2.033333	8026	
2	AT	2006	3	2.739429	3.221367	5.380802	5.7	1.800000	8027	
3	AT	2006	4	2.739429	3.594467	5.380802	5.7	1.433333	8028	
4	AT	2007	1	2.242467	3.820333	6.119790	5.3	1.766667	8029	

Next steps: [Generate code with df](#) [New interactive sheet](#)

```
df_panel = df.set_index(["country","time_id"]).sort_index()
df_panel.head()
```

	country	time_id	year	quarter	npl_ratio	euribor_3m	gdp_growth	unemployment	inflation_yoy	grid icon
AT	8025	2006	1	1	2.739429	2.611567	5.380802	5.7	1.433333	info icon
	8026	2006	2	2	2.739429	2.889500	5.380802	5.7	2.033333	
	8027	2006	3	3	2.739429	3.221367	5.380802	5.7	1.800000	
	8028	2006	4	4	2.739429	3.594467	5.380802	5.7	1.433333	
	8029	2007	1	1	2.242467	3.820333	6.119790	5.3	1.766667	

Next steps: [Generate code with df_panel](#) [New interactive sheet](#)

```
vars_analysis = ["npl_ratio","gdp_growth","unemployment","inflation_yoy","euribor_3m"]
print(df_panel[vars_analysis].describe())
```

	npl_ratio	gdp_growth	unemployment	inflation_yoy	euribor_3m
count	1288.000000	1288.000000	1288.000000	1288.000000	1288.000000
mean	6.262905	4.332987	8.822360	2.477666	0.964973
std	8.042270	5.759178	4.539123	3.191348	1.636366
min	0.145588	-21.188764	3.100000	-3.866667	-0.566400
25%	1.992508	1.906078	5.900000	0.633333	-0.319500
50%	3.351983	4.012819	7.350000	1.816667	0.223500
75%	6.375134	7.305004	10.300000	3.166667	1.772133
max	47.747846	35.737719	27.800000	24.166667	4.981800

	npl_ratio	gdp_growth	unemployment	inflation_yoy	euribor_3m
npl_ratio	1.000000	-0.235386	0.624301	-0.261648	-0.199448
gdp_growth	-0.235386	1.000000	-0.340156	0.384545	0.099805
unemployment	0.624301	-0.340156	1.000000	-0.265070	-0.192088
inflation_yoy	-0.261648	0.384545	-0.265070	1.000000	0.292140
euribor_3m	-0.199448	0.099805	-0.192088	0.292140	1.000000

```
from linearmodels.panel import PanelOLS
import statsmodels.api as sm

y = df_panel["npl_ratio"]
X = df_panel[["gdp_growth","unemployment","inflation_yoy","euribor_3m"]]

fe = PanelOLS(y, X, entity_effects=True, check_rank=False)
fe_res = fe.fit(cov_type="clustered", cluster_entity=True)
print(fe_res)
```

PanelOLS Estimation Summary					
Dep. Variable:	npl_ratio	R-squared:	0.5060		
Estimator:	PanelOLS	R-squared (Between):	0.1981		
No. Observations:	1288	R-squared (Within):	0.5060		
Date:	Tue, Nov 25 2025	R-squared (Overall):	0.3497		
Time:	22:12:57	Log-likelihood	-3682.2		
Cov. Estimator:	Clustered	F-statistic:	323.96		
Entities:	19	P-value	0.0000		
Avg Obs:	67.789	Distribution:	F(4,1265)		
Min Obs:	12.000				
Max Obs:	80.000	F-statistic (robust):	8.2919		
		P-value	0.0000		
Time periods:	80	Distribution:	F(4,1265)		
Avg Obs:	16.100				
Min Obs:	4.0000				
Max Obs:	19.000				

Parameter Estimates						
Parameter	Std. Err.	T-stat	P-value	Lower CI	Upper CI	
gdp_growth	0.0231	0.0713	0.3236	0.7463	-0.1168	0.1630
unemployment	1.3405	0.2703	4.9592	0.0000	0.8102	1.8707
inflation_yoy	-0.0331	0.0793	-0.4174	0.6765	-0.1888	0.1225
euribor_3m	-0.3762	0.1661	-2.2651	0.0237	-0.7020	-0.0504

F-test for Poolability: 81.939
P-value: 0.0000
Distribution: F(18,1265)

Included effects: Entity

```
from linearmodels.panel import RandomEffects

X_re = sm.add_constant(X)
re = RandomEffects(y, X_re)
re_res = re.fit()
print(re_res)
```

RandomEffects Estimation Summary

Dep. Variable:	npl_ratio	R-squared:	0.5028
Estimator:	RandomEffects	R-squared (Between):	0.1520
No. Observations:	1288	R-squared (Within):	0.5060
Date:	Tue, Nov 25 2025	R-squared (Overall):	0.3780
Time:	22:13:17	Log-likelihood	-3691.9
Cov. Estimator:	Unadjusted	F-statistic:	324.35
Entities:	19	P-value	0.0000
Avg Obs:	67.789	Distribution:	F(4,1283)
Min Obs:	12.000		
Max Obs:	80.000	F-statistic (robust):	324.54
		P-value	0.0000
Time periods:	80	Distribution:	F(4,1283)
Avg Obs:	16.100		
Min Obs:	4.0000		
Max Obs:	19.000		

Parameter Estimates

Parameter	Std. Err.	T-stat	P-value	Lower CI	Upper CI	
const	-5.1863	1.2296	-4.2179	0.0000	-7.5986	-2.7741
gdp_growth	0.0221	0.0247	0.8935	0.3718	-0.0264	0.0706
unemployment	1.3303	0.0431	30.855	0.0000	1.2457	1.4149
inflation_yoy	-0.0363	0.0447	-0.8118	0.4170	-0.1239	0.0514
euribor_3m	-0.3759	0.0778	-4.8301	0.0000	-0.5286	-0.2232

```
from scipy import stats

beta = ["gdp_growth", "unemployment", "inflation_yoy", "euribor_3m"]

b_fe = fe_res.params[beta]
b_re = re_res.params[beta]

V_fe = fe_res.cov.loc[beta, beta]
V_re = re_res.cov.loc[beta, beta]

diff = b_fe - b_re
Vdiff = V_fe - V_re

H = float(diff.T @ np.linalg.inv(Vdiff) @ diff)
p = 1 - stats.chi2.cdf(H, len(diff))

print("Hausman chi2:", H)
print("p-value:", p)

Hausman chi2: -0.03990132049242252
p-value: 1.0
```

```
df_panel.reset_index().to_csv("macro_clean_final.csv", index=False)
from google.colab import files
files.download("macro_clean_final.csv")
```



